

Chapter 25. Boost.PropertyTree

With the class `boost::property_tree::ptree`, [Boost.PropertyTree](#) provides a tree structure to store key/value pairs. Tree structure means that a trunk exists with numerous branches that have numerous twigs. A file system is a good example of a tree structure. File systems have a root directory with subdirectories that themselves can have subdirectories and so on.

To use `boost::property_tree::ptree`, include the header file `boost/property_tree/ptree.hpp`. This is a master header file, so no other header files need to be included for `Boost.PropertyTree`.

Example 25.1. Accessing data in `boost::property_tree::ptree`

```
#include <boost/property_tree/ptree.hpp>
#include <iostream>

using boost::property_tree::ptree;

int main()
{
    ptree pt;
    pt.put("C:.Windows.System", "20 files");

    ptree &c = pt.get_child("C:");
    ptree &windows = c.get_child("Windows");
    ptree &system = windows.get_child("System");
    std::cout << system.get_value<std::string>() << '\n';
}
```

[Example 25.1](#) uses `boost::property_tree::ptree` to store a path to a directory. This is done with a call to `put()`. This member function expects two parameters because `boost::property_tree::ptree` is a tree structure that saves key/value pairs. The tree doesn't just consist of branches and twigs, a value must be assigned to each branch and twig. In [Example 25.1](#) the value is “20 files”.

The first parameter passed to `put()` is more interesting. It is a path to a directory. However, it doesn't use the backslash, which is the common path separator on Windows. It uses the dot.

You need to use the dot because it's the separator `Boost.PropertyTree` expects for keys. The parameter “C:.Windows.System” tells `pt` to create a branch called C: with a branch called Windows that has another branch called System. The dot creates the nested structure of branches. If “C:\Windows\System” had been passed as the parameter, `pt` would only have one branch called C:\Windows\System.

After the call to `put()`, `pt` is accessed to read the stored value “20 files” and write it to standard output. This is done by jumping from branch to branch – or directory to directory.

To access a subbranch, you call `get_child()`, which returns a reference to an object of the same type `get_child()` was called on. In [Example 25.1](#), this is a reference to `boost::property_tree::ptree`. Because every branch can have subbranches, and because there is no structural difference between higher and lower branches, the same type is used.

The third call to `get_child()` retrieves the `boost::property_tree::ptree`, which represents the directory System. `get_value()` is called to read the value that was stored at the beginning of the example with `put()`.

Please note that `get_value()` is a function template. You pass the type of the return value as a template parameter. That way `get_value()` can do an automatic type conversion.

Example 25.2. Accessing data in `basic_ptree<std::string, int>`

```
#include <boost/property_tree/ptree.hpp>
#include <utility>
#include <iostream>

int main()
{
    typedef boost::property_tree::basic_ptree<std::string, int> ptree;
    ptree pt;
    pt.put(ptree::path_type{"C:\\Windows\\System", '\\'}, 20);
    pt.put(ptree::path_type{"C:\\Windows\\Cursors", '\\'}, 50);

    ptree &windows = pt.get_child(ptree::path_type{"C:\\Windows", '\\'});
    int files = 0;
    for (const std::pair<std::string, ptree> &p : windows)
        files += p.second.get_value<int>();
    std::cout << files << '\n';
}
```

There are two changes in [Example 25.2](#) compared with [Example 25.1](#). These changes are to save paths to directories and the number of files in directories more easily. First, paths use a backslash as the separator when passed to `put()`. Secondly, the number of files is stored as an `int`.

By default, Boost.PropertyTree uses a dot as the separator for keys. If you need to use another character, such as the backslash, as the separator, you don't pass the key as a string to `put()`. Instead you wrap it in an object of type `boost::property_tree::ptree::path_type`. The constructor of this class, which depends on `boost::property_tree::ptree`, takes the key as its first parameter and the separator character as its second parameter. That way, you can use a path such as C:\Windows\System, as shown in [Example 25.2](#), without having to replace backslashes with dots.

`boost::property_tree::ptree` is based on the class template `boost::property_tree::basic_ptree`. Because keys and values are often strings, `boost::property_tree::ptree` is predefined. However, you can use `boost::property_tree::basic_ptree` with different types for keys and values. The tree in [Example 25.2](#) uses an `int` to store the number of files in a directory rather than a string.

`boost::property_tree::ptree` provides the member functions `begin()` and `end()`. However, `boost::property_tree::ptree` only lets you iterate over the branches in one level. [Example 25.2](#) iterates over the subdirectories of C:\Windows. You can't get an iterator to iterate over all branches in all levels.

The `for` loop in [Example 25.2](#) reads the number of files in all subdirectories of C:\Windows to calculate a total. As a result, the example displays **70**. The example doesn't access objects of

type `ptree` directly. Instead it iterates over elements of type `std::pair<std::string, ptree>`. **first** contains the key of the current branch. That is System and Cursors in [Example 25.2](#). **second** provides access to an object of type `ptree`, which represents the possible subdirectories. In the example, only the values assigned to System and Cursors are read. As in [Example 25.1](#), the member function `get_value()` is called.

`boost::property_tree::ptree` only stores the value of the current branch, not its key. You can get the value with `get_value()`, but there is no member function to get the key. The key is stored in `boost::property_tree::ptree` one level up. This also explains why the `for` loop iterates over elements of type `std::pair<std::string, ptree>`.

Example 25.3. Accessing data with a translator

```
#include <boost/property_tree/ptree.hpp>
#include <boost/optional.hpp>
#include <iostream>
#include <cstdlib>

struct string_to_int_translator
{
    typedef std::string internal_type;
    typedef int external_type;

    boost::optional<int> get_value(const std::string &s)
    {
        char *c;
        long l = std::strtol(s.c_str(), &c, 10);
        return boost::make_optional(c != s.c_str(), static_cast<int>(l));
    }
};

int main()
{
    typedef boost::property_tree::iptree ptree;
    ptree pt;
    pt.put(ptree::path_type{"C:\\Windows\\System", '\\\\'}, "20 files");
    pt.put(ptree::path_type{"C:\\Windows\\Cursors", '\\\\'}, "50 files");

    string_to_int_translator tr;
    int files =
        pt.get<int>(ptree::path_type{"c:\\windows\\system", '\\\\'}, tr) +
        pt.get<int>(ptree::path_type{"c:\\windows\\cursors", '\\\\'}, tr);
    std::cout << files << '\n';
}
```

[Example 25.3](#) uses with `boost::property_tree::iptree` another predefined tree from Boost.PropertyTree. In general, this type behaves like `boost::property_tree::ptree`. The only difference is that `boost::property_tree::iptree` doesn't distinguish between lower and upper case. For example, a value stored with the key C:\Windows\System can be read with c:\windows\system.

Unlike [Example 25.1](#), `get_child()` isn't called multiple times to access subbranches. Just as `put()` can be used to store a value in a subbranch directly, a value from a subbranch can be read with `get()`. The key is defined the same way – for example using `boost::property_tree::iptree::path_type`.

Like `get_value()`, `get()` is a function template. You have to pass the type of the return value as a template parameter. `Boost.PropertyTree` does an automatic type conversion.

To convert types, `Boost.PropertyTree` uses *translators*. The library provides a few translators out of the box that are based on streams and can convert types automatically.

[Example 25.3](#) defines the translator `string_to_int_translator`, which converts a value of type `std::string` to `int`. The translator is passed as an additional parameter to `get()`. Because the translator is just used to read, it only defines one member function, `get_value()`. If you want to use the translator for writing, too, then you would need to define a member function `put_value()` and then pass the translator as an additional parameter to `put()`.

`get_value()` returns a value of the type that is used in `pt`. However, because a type conversion doesn't always succeed, `boost::optional` is used. If a value is stored in [Example 25.3](#) that can't be converted to an `int` with `std::strtol()`, an empty object of type `boost::optional` will be returned.

Please note that a translator must also define the two types `internal_type` and `external_type`. If you need to convert types when storing data, define `put_value()` similar to `get_value()`.

If you modify [Example 25.3](#) to store the value "20" instead of value "20 files", `get_value()` can be called without passing a translator. The translators provided by `Boost.PropertyTree` can convert from `std::string` to `int`. However, the type conversion only succeeds when the entire string can be converted. The string must not contain any letters. Because `std::strtol()` can do a type conversion as long as the string starts with digits, the more liberal translator `string_to_int_translator` is used in [Example 25.3](#).

Example 25.4. Various member functions of `boost::property_tree::ptree`

```
#include <boost/property_tree/ptree.hpp>
#include <utility>
#include <iostream>

using boost::property_tree::ptree;

int main()
{
    ptree pt;
    pt.put("C:.Windows.System", "20 files");

    boost::optional<std::string> c = pt.get_optional<std::string>("C:");
    std::cout << std::boolalpha << c.is_initialized() << '\n';

    pt.put_child("D:.Program Files", ptree{"50 files"});
    pt.add_child("D:.Program Files", ptree{"60 files"});

    ptree d = pt.get_child("D:");
    for (const std::pair<std::string, ptree> &p : d)
        std::cout << p.second.get_value<std::string>() << '\n';

    boost::optional<ptree> e = pt.get_child_optional("E:");
    std::cout << e.is_initialized() << '\n';
}
```

You can call the member function `get_optional()` if you want to read the value of a key, but you aren't sure if the key exists. `get_optional()` returns the value in an object of type `boost::optional`. The object is empty if the key wasn't found. Otherwise, `get_optional()` works the same as `get()`.

It might seem like `put_child()` and `add_child()` are the same as `put()`. The difference is that `put()` creates only a key/value pair while `put_child()` and `add_child()` insert an entire subtree. Note that an object of type `boost::property_tree::ptree` is passed as the second parameter to `put_child()` and `add_child()`.

The difference between `put_child()` and `add_child()` is that `put_child()` accesses a key if that key already exists, while `add_child()` always inserts a new key into the tree. That's why the tree in [Example 25.4](#) has two keys called "D:.Program Files". Depending on the use case, this can be confusing. If a tree represents a file system, there shouldn't be two identical paths. You have to avoid inserting identical keys if you don't want duplicates in a tree.

[Example 25.4](#) displays the value of the keys below "D:" in the `for` loop. The example writes `50 files` and `60 files` to standard output, which proves there are two identical keys called "D:.Program Files".

The last member function introduced in [Example 25.4](#) is `get_child_optional()`. This function is used like `get_child()`. `get_child_optional()` returns an object of type `boost::optional`. You call `boost::optional` if you aren't sure whether a key exists.

Example 25.5. Serializing a `boost::property_tree::ptree` in the JSON format

```
#include <boost/property_tree/ptree.hpp>
#include <boost/property_tree/json_parser.hpp>
#include <iostream>

using namespace boost::property_tree;

int main()
{
    ptree pt;
    pt.put("C:.Windows.System", "20 files");
    pt.put("C:.Windows.Cursors", "50 files");

    json_parser::write_json("file.json", pt);

    ptree pt2;
    json_parser::read_json("file.json", pt2);

    std::cout << std::boolalpha << (pt == pt2) << '\n';
}
```

Boost.PropertyTree does more than just provide structures to manage data in memory. As can be seen in [Example 25.5](#), the library also provides functions to save a `boost::property_tree::ptree` in a file and load it from a file.

The header file `boost/property_tree/json_parser.hpp` provides access to the functions `boost::property_tree::json_parser::write_json()` and `boost::property_tree::json_parser::read_json()`. These functions make it possible to

save and load a `boost::property_tree::ptree` serialized in the JSON format. That way you can support configuration files in the JSON format.

If you want to call functions that store a `boost::property_tree::ptree` in a file or load it from a file, you must include header files such as `boost/property_tree/json_parser.hpp`. It isn't sufficient to only include `boost/property_tree/ptree.hpp`.

In addition to the functions `boost::property_tree::json_parser::write_json()` and `boost::property_tree::json_parser::read_json()`, Boost.PropertyTree provides functions for additional data formats. You use

`boost::property_tree::ini_parser::write_ini()` and `boost::property_tree::ini_parser::read_ini()` from `boost/property_tree/ini_parser.hpp` to support INI-files. With `boost::property_tree::xml_parser::write_xml()` and `boost::property_tree::xml_parser::read_xml()` from `boost/property_tree/xml_parser.hpp`, data can be loaded and stored in XML format. With `boost::property_tree::info_parser::write_info()` and `boost::property_tree::info_parser::read_info()` from `boost/property_tree/info_parser.hpp`, you can access another format that was developed and optimized to serialize trees from Boost.PropertyTree.

None of the supported formats guarantees that a `boost::property_tree::ptree` will look the same after it has been saved and reloaded. For example, the JSON format can lose type information because `boost::property_tree::ptree` can't distinguish between `true` and "true". The type is always the same. Even if the various functions make it easy to save and load a `boost::property_tree::ptree`, don't forget that Boost.PropertyTree doesn't support the formats completely. The main focus of the library is on the structure `boost::property_tree::ptree` and not on supporting various data formats.

Exercise

Create a program that loads this JSON-file and writes the names of all animals to standard output. If "all" is set to `true` the program should not only write the names but all properties of all animals to standard output:

```
{
  "animals": [
    {
      "name": "cat",
      "legs": 4,
      "has_tail": true
    },
    {
      "name": "spider",
      "legs": 8,
      "has_tail": false
    }
  ],
  "log": {
    "all": true
  }
}
```

Solutions

theboostcplibraries.com



Solutions from
the expert to all
exercises in the
book for \$9.99